



Course FSD311: Generative Adversarial Networks (GANs)

Thomas Pierrot: t.pierrot@instadeep.com

Contents

Generative Models

Variational Autoencoders (VAEs)

Generative Adversarial Networks (GANs)

Transpose Convolution and DCGANs

Conditional GANs

Motivation

Yann Lecun (Director of Facebook AI Research Paris and Professor at NYU) on Twitter:

"There are many interesting recent development in deep learning, probably too many for me to describe them all here. But there are a few ideas that caught my attention enough for me to get personally involved in research projects. **The most important one, in my opinion, is adversarial training** (also called GAN for Generative Adversarial Networks). This is an idea that was originally proposed by Ian Goodfellow when he was a student with Yoshua Bengio at the University of Montreal (he since moved to Google Brain and recently to OpenAI). **This, and the variations that are now being proposed is the most interesting idea in the last years in ML, in my opinion.**"

Motivation

Are those pictures real ?



Motivation

Are those pictures real ?



No ! They've been dreamed by a neural network.

Source : Progressive growing of GANs paper

Generative Models


- **Goal** : Generate with a model (in this case a NN) data that look real but does not already exist.
- **Example** : Generate pictures that look like pictures from a given dataset.
- **Problem** : No natural input in this situation.
- **Solution**: Random numbers as inputs.
- **Finally** : Model that maps a given distribution (ex: gaussian) towards our data distribution.

Generative Models


Real data (used to train the model)



Random noise



Generative model
(or Generator)



Data that looks real



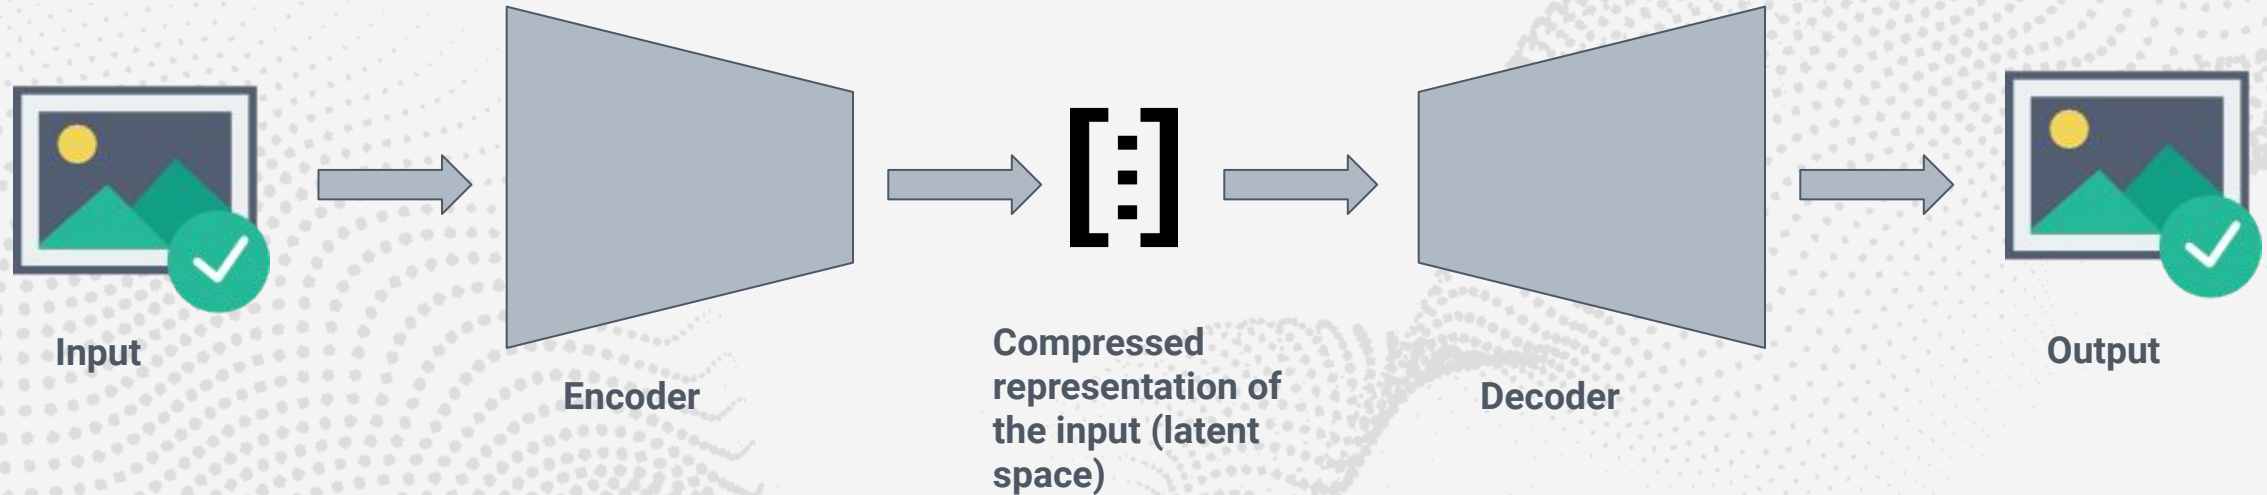
Generative Models

- **Question : How to define a meaningful loss to train the model ?**
- **Possible solution : Minimize the distance between the generator output and its closest neighbour in the dataset.**
- **Issue : Extremely expensive and leads to poor results.**
- **Solution : Use a second network to train the generator.**
 - **First possibility: VAE: Variational autoencoders**
 - **Second possibility: GANs: Generative Adversarial Networks**

Variational Autoencoders (VAE)

Autoencoders

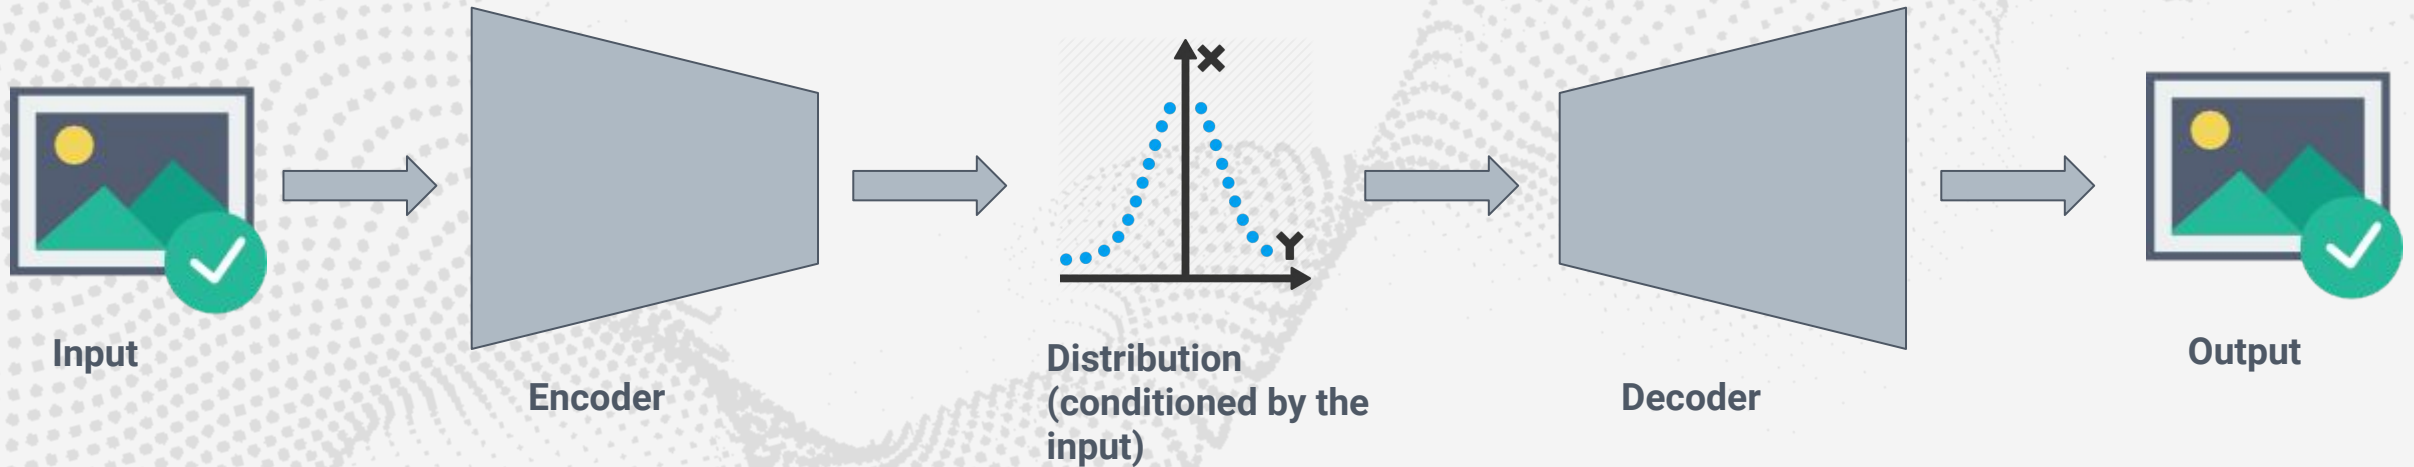
The autoencoder is trained to return an output that equals its input. Thus, we minimize a distance between the network output and its input.



Autoencoder may be used for compression, to learn representations, but also as generative models.

Variational Autoencoders

Instead of learning deterministic representations, VAEs encode its input as a distribution. Then the decoder is trained to decode any vector sampled according to this distribution.



In practice, we assume the distribution to be a gaussian and the encoder network predicts a mean vector and a covariance matrix.

Let's write it mathematically !

Notations

- **Inputs :** $\mathbf{x} \in \mathcal{X}$
- **Encoder :** $q_{\theta} : x \in \mathcal{X} \rightarrow q_{\theta}(z|x)$
- **Decoder :** $p_{\phi} : z \sim q_{\theta}(z|x) \rightarrow x \in \mathcal{X}$

Atomic loss (for one input):
$$l_i(\theta, \phi) = \underbrace{-\mathbb{E}_{z \sim q_{\theta}(z|x_i)}[\log(p_{\phi}(x_i|z))]}_{\text{Reconstruction loss}} + \underbrace{\text{KL}(q_{\theta}(z|x_i) || p(z))}_{\text{Regularization loss}}$$


Reconstruction loss

Regularization loss

$$p(z) \sim \mathcal{N}(0, 1)$$

VAE as a generative model

Once trained, the decoder can be used alone as a generative model. To use it, vectors are sampled according to a normal gaussian and decoded into data that looks real.



Major drawback: the reconstruction loss. When the input is a picture, considering a reconstruction loss on its pixels is not really meaningful and leads to poor results.

Generative Adversarial Networks (GANs)

GAN fundamental idea

- Defining meaningful distances between data (especially images) might very difficult and often require expert knowledge.
- In a GAN this distance, used to train the Generator, is learned implicitly through a second neural network called Discriminator.

Discriminator network principle

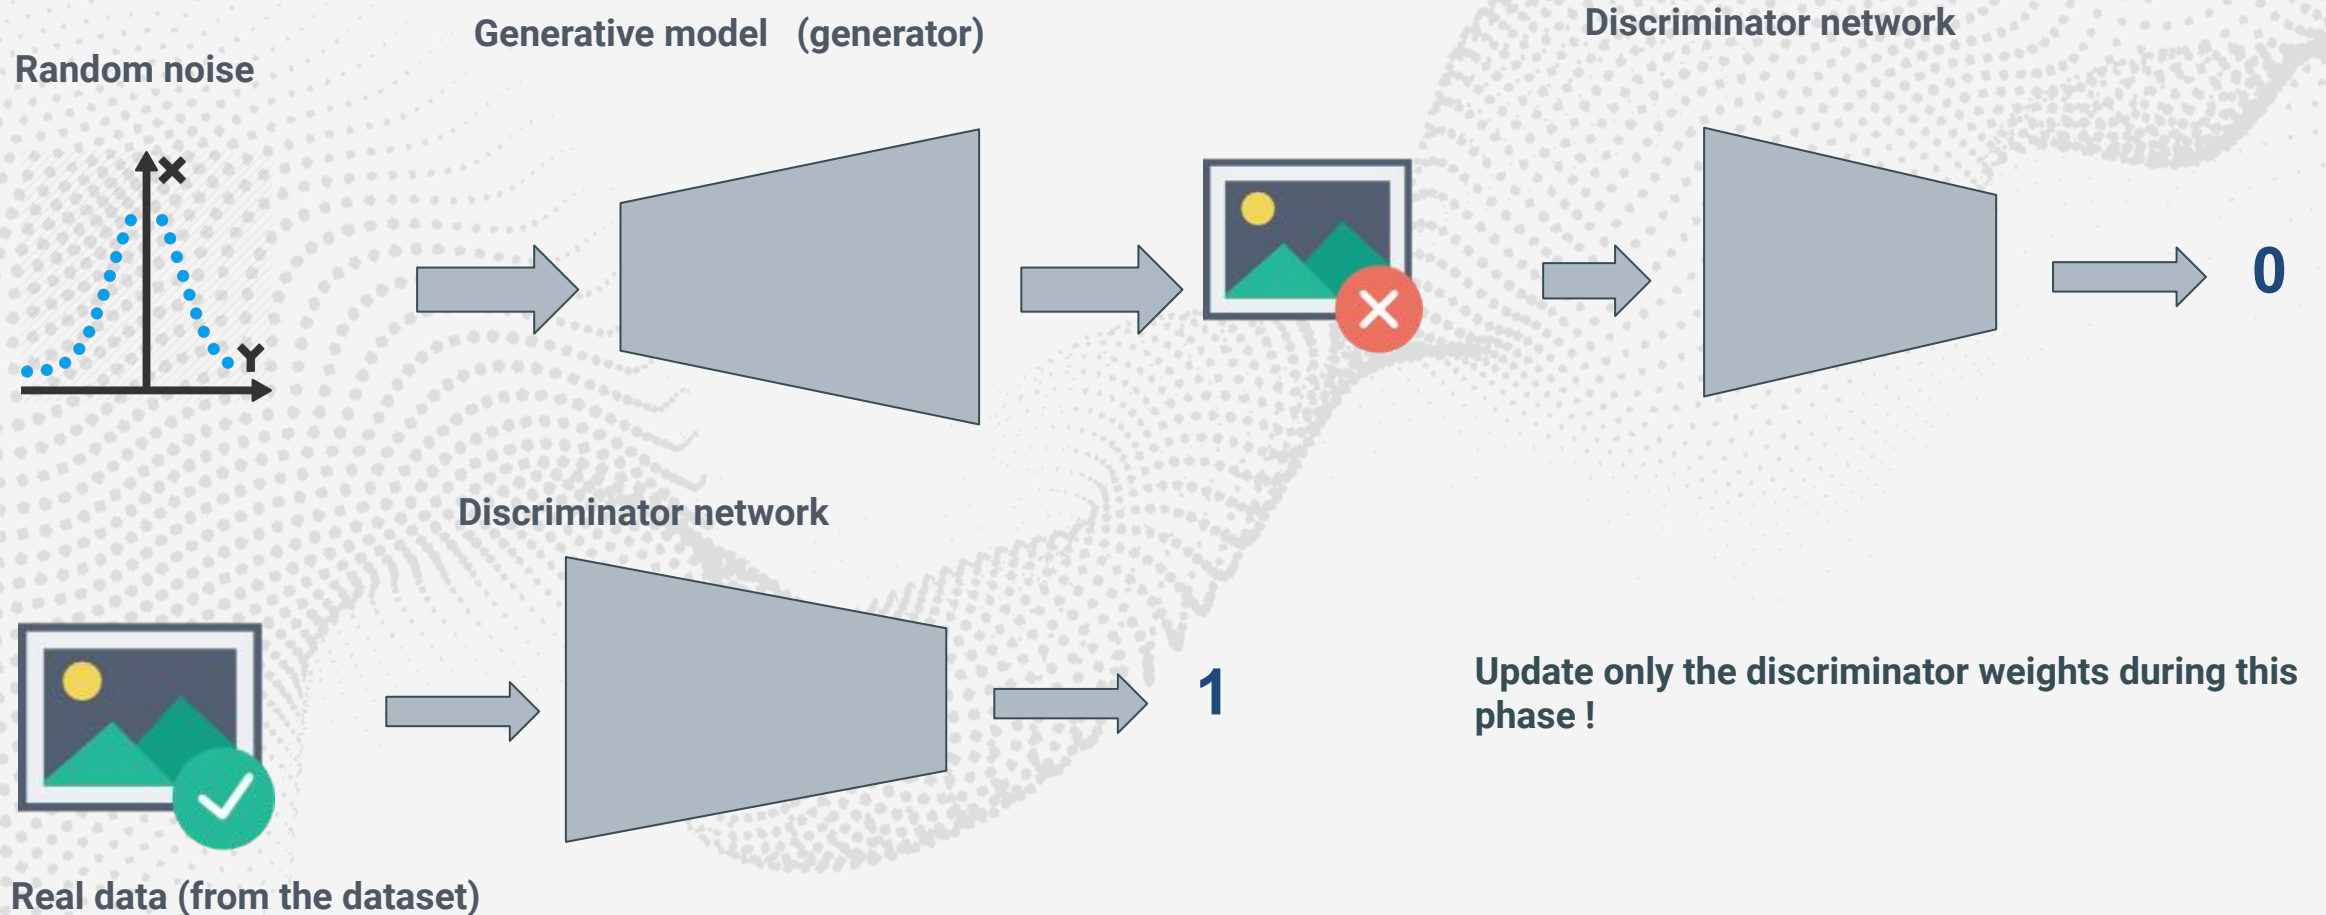


The discriminator returns the probability for a given input to be real.


Adversarial training

- Train the discriminator to make the difference between fake and real data (standard supervised training).
- Train the generator to fool the discriminator.
- Repeat it until the discriminator is completely lost (returns always 0.5).
- At this point the generator returns data that looks real.

Train the discriminator



Train the generator



Update only the generator weights during this phase !

Let's write it mathematically !

Notations

- **Inputs :** $\mathbf{x} \in \mathcal{X}$
- **Random noise :** $\mathbf{z} \in \mathcal{G}$, example : $\mathbf{z} \sim \mathcal{N}(0, 1)$
- **Generator :** $G : \mathcal{G} \longrightarrow \mathcal{X}$
- **Discriminator :** $D : \mathcal{X} \longrightarrow \mathbb{R}$

How it works

- Discriminator loss:

$$L_D = - \sum_{i=1}^n \underbrace{\log(1 - D(G(z_i)))}_{L_D^{false}} + \underbrace{\log(D(x_i))}_{L_D^{true}}$$

- Generator loss:

$$L_G = \sum_i^n \log(1 - D(G(z_i)))$$

Warning: I took the convention where the losses are to be **minimized**.


Note: As usual, the full sum is replaced with a sub-sum: we use mini-batches to compute the gradients.

Practical algorithm

- **Initialize both networks**
- **Training:**
 - Sample a mini-batch of noise vectors and sample a mini-batch of input data.
 - Update the discriminator weights with those batches (gradient descent on its loss).
 - Sample a mini-batch of noise vectors.
 - Update the generator weights with this batch (gradient descent as well).
- **Repeat training phase until convergence**

Architectures used in the original paper


The discriminator (dense NN)



Takes a flattened image and returns a scalar value

Architectures used in the original paper

The generator (dense NN)




Takes a vector of random number and returns a flattened image

Toward the DCGANs

- Now we want to work with large images.
- Problem: feed directly a large image into a dense NN will make the number of weights blow.
- Standard idea: use a convolutional NN instead of a fully dense NN.

Reminder : convNets


- Main idea : reduce input size (extract its features) before feeding it into a standard (dense) NN.
- Historically : hand-made feature engineering.
- Since 2012 : automatic feature extraction with convolution layers.



Example of convolution with a filter of size 3x3 on a 4x4 image. Result of the a convolution: a 2x2 image.

Example for the discriminator

- Extract the input image features through several convolutional layers.
- Use the extracted features inside a NN layer to determine whether the image is real or fake.




A discriminator example on MNIST

What about the generator ?

- Convolution decreases the data size: it is perfect for the discriminator.
- However in the generator we need an operation that increases the data size !
- One might use standard interpolation techniques.
- But it leads to poor results (they are not really “learnable”).
- Answer: the transposed convolution (also called deconvolution or up-sampling).

Toward the transposed-convolution

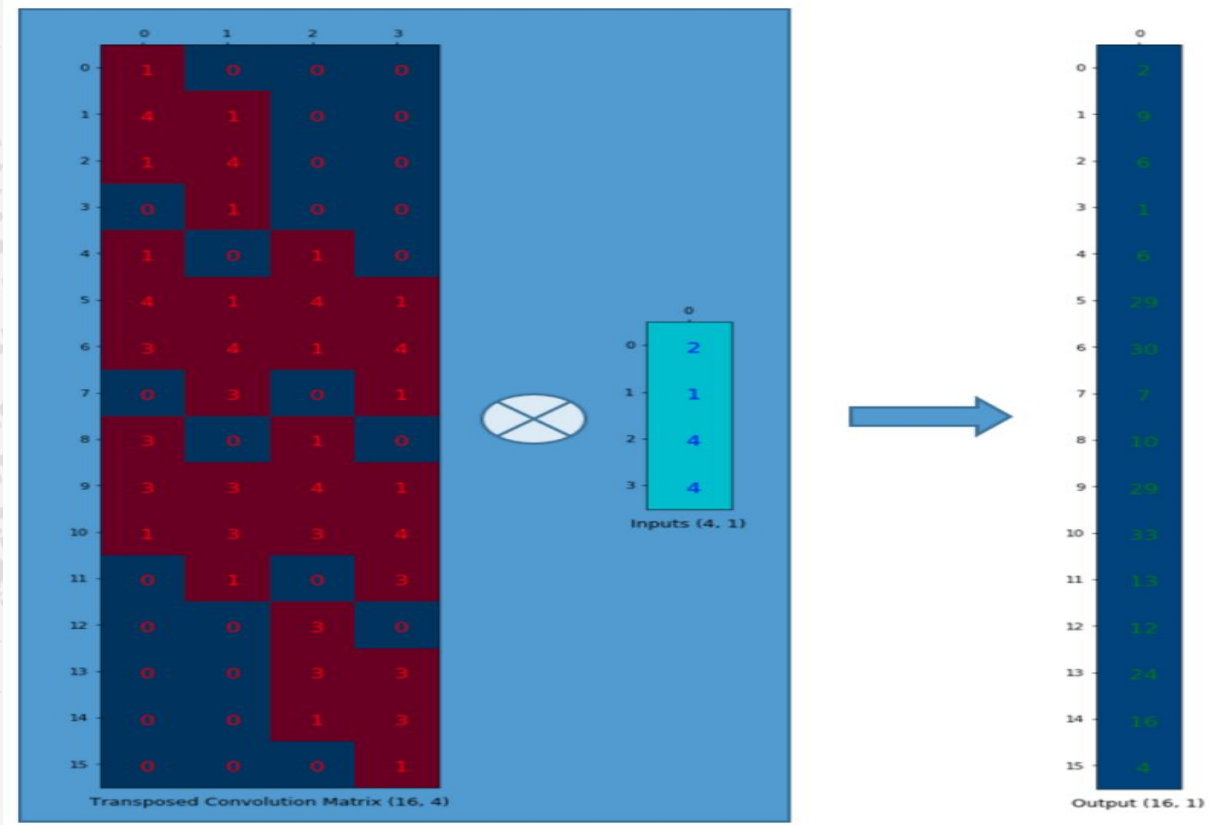
Convolution as a matrix vector product operation.



Example: convolution over a 4x4 image (flattened as a vector of size 16) with a 3x3 filter. It gives a 2x2 image

Toward the transposed-convolution

We define the inverse operation: matrix vector product with the transposed convolution matrix.




As for the convolution, the matrix coefficients are learnt!

Example: deconvolution over a 2x2 image (flattened as a vector of size 4) with a 3x3 filter. It gives a 4x4 image

Example for the generator

- Perform up-sampling operations to transform the input noise vector into an image.



A generator example on MNIST

DCGAN

- Radford *et al.* (2015) present the Deep Convolutional GAN (DCGAN).



Figure 3: Generated bedrooms after five epochs of training. There appears to be evidence of visual under-fitting via repeated noise textures across multiple samples such as the base boards of some of the beds.

- They generate fake pictures of bedrooms to demonstrate their architecture efficiency.

SRGAN

- Ledig *et al.* (2016): **super-resolution GAN model, SRGAN that can take a low resolution photo and generate a high-resolution sample.**




Figure 2: From left to right: bicubic interpolation, deep residual network optimized for MSE, deep residual generative adversarial network optimized for a loss more sensitive to human perception, original HR image. Corresponding PSNR and SSIM are shown in brackets. [4× upscaling]

- **This problem is classical supervised learning. Here, the GAN scheme is used as an alternative loss to the standard RMS (or L2) loss.**

Conditional generative models


- GANs drawback: it generates data without “control on the output”.
- Example: if it learns to generate MNIST digits, it will generate any digit without distinction.
- One might desire to generate (with the same network) a given digit (a given label).
- Answer: we can condition the network output.
- In this case, we want to generate data given a condition c .
- Example: generate a MNIST digit given its label.

Conditional GANs

- Adapt GAN to add a condition on the input: **CGANs** (M. Mirza et al., 2014)
- Simple idea: add an input to both the generator and the discriminator.
- The generator takes a noise vector and a condition (example a label) and generates conditioned (by this label) image.
- The discriminator takes an image and the condition and returns the conditioned (by this label) probability that this image is real.
- Note: most of the time the condition is a label, but it can be a description or other type of data.

CGAN architecture


The discriminator (dense NN)



Takes a flattened image and a condition and returns a scalar value.

CGAN architecture

The generator (dense NN)



Takes a vector of random number and a condition and returns a flattened image.

Adapt GAN losses to train the CGAN

- Discriminator loss:

$$L_D = - \sum_i^n \log (1 - D(G(z_i|c_i)|c_i)) + \log (D(x_i|c_i))$$

- Generator loss:

$$L_G = \sum_i^n \log (1 - D(G(z_i|c_i)))$$

Warning: To train the discriminator **the same label** must be fed into the generator and the discriminator.

CGAN Practical algorithm

- **Initialize both networks**
- **Training:**
 - Sample a mini-batch of noise vectors, sample a mini-batch of input data and a batch of labels (c_i).
 - Update the discriminator weights with those batches (gradient descent on its loss).
 - Sample a mini-batch of noise vectors and a batch of labels.
 - Update the generator weights with this batch (gradient descent as well).
- **Repeat training phase until convergence**

BigGan: DeepMind 2018

- Of course, CGAN and DCGAN may be mixed !



Figure 1: Class-conditional samples generated by our model.

- This year deepMind (A. Brock et al., 2018) proposed an algorithm BigGan that learns the conditional probability of the imageNet dataset (1000 classes !)

State of the Art: 3D

- GAN models can generate 3D objects, 3D environments and even VR environments.
- Wu *et al.* in a ground-breaking article recently published (Jan 2017) showed that it's possible to generate 3D furniture out of simple photos!



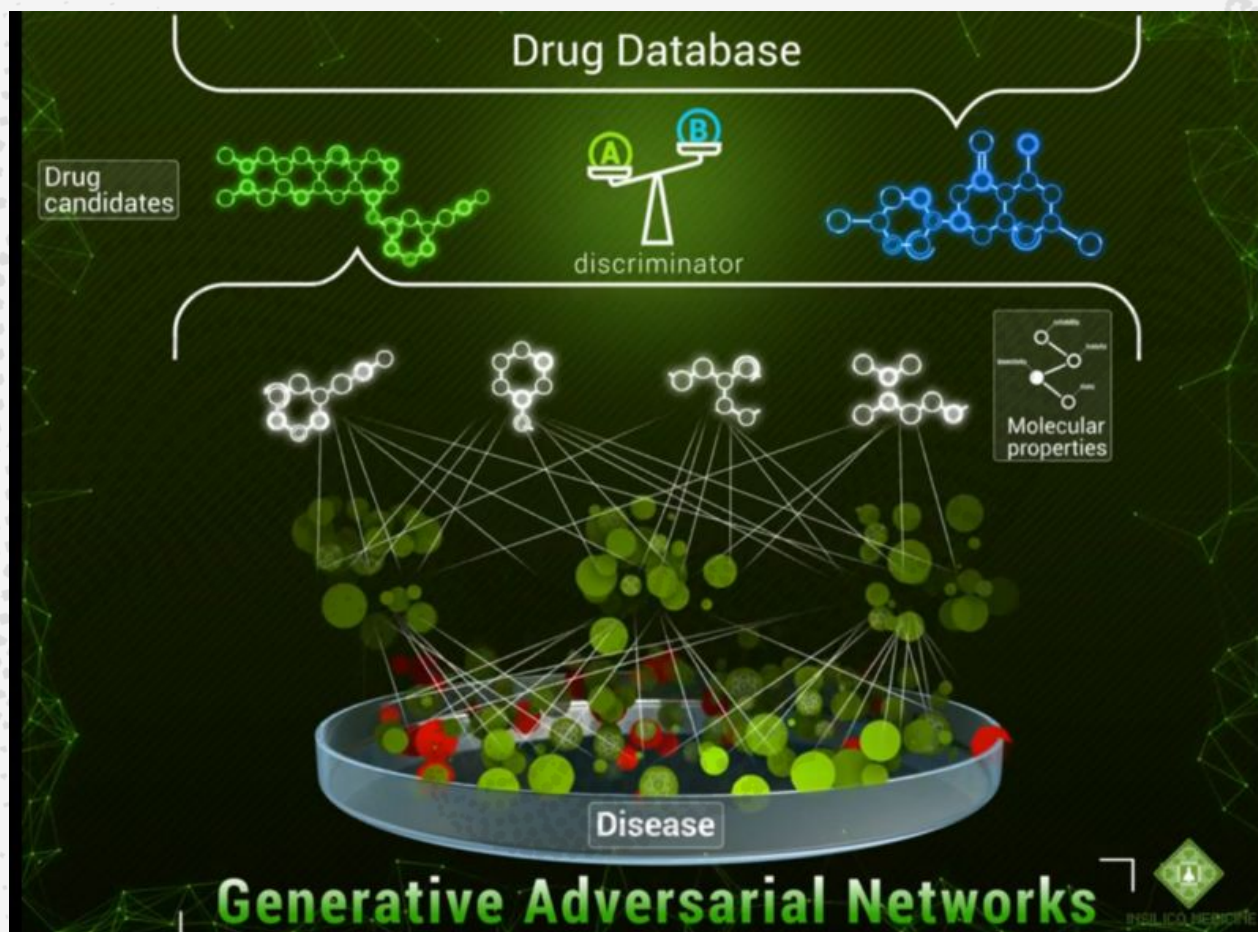
- Even more incredible you can even do arithmetic on these 3D models to create **new, never seen before furniture!**



- More state of the art examples: protein research, astronomy

Other applications: drug discovery

See: <https://medium.com/neuromation-io-blog/creating-molecules-from-scratch-i-drug-discovery-with-generative-adversarial-networks-9d42cc496fc6>




Tutorials

- First part: download, read and complete the notebook “GAN-fill the gaps” that is available on GitHub.
- Second part: take this notebook and adapt it to train a CGAN on MNIST.



Examples of pictures generated by the GAN (we don't control the digits generated).



Examples of pictures generated by the CGAN (we control the digits generated).